

Large Scale Computing

Lab 6

Adrian Madej

1. Short description of running the application

As part of the project, a web application was deployed in a local Kubernetes cluster using kind. The application includes the following components:

- An NFS server with a dynamic provisioner, installed via Helm,
- A PersistentVolumeClaim with ReadWriteMany access mode,
- A Pod running an nginx HTTP server that mounts the NFS volume,
- A Kubernetes Job that writes a sample index.html file to the shared volume,
- A NodePort Service that exposes the HTTP server for external access.

The entire configuration and commands used to deploy the application are included in the GitHub repository - [Large-Scale-Computing/lab_6](https://github.com/ltmollo/Large-Scale-Computing) at main · ltmollo/Large-Scale-Computing .

A screenshot of a web page:



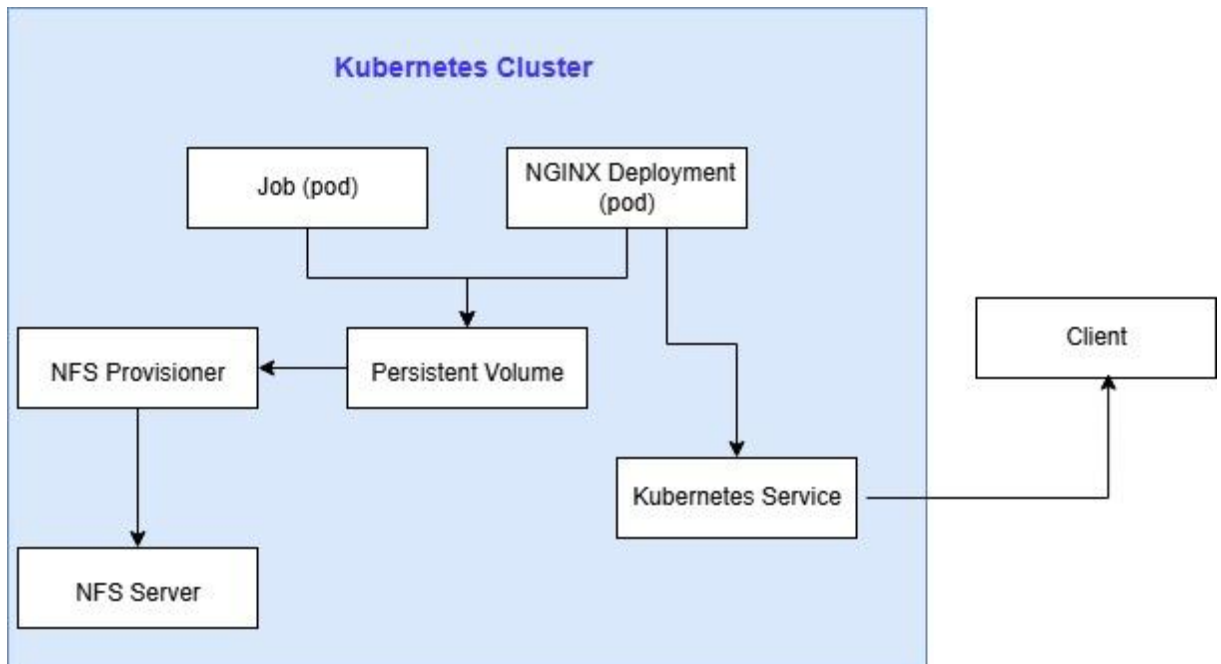
Example outputs:

```
PS C:\Users\ltmol\Desktop\studia\lsc\lab_6> kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
nfs-server-nfs-server-provisioner-0 1/1     Running   0           73s
PS C:\Users\ltmol\Desktop\studia\lsc\lab_6> kubectl get storageclass
NAME                                PROVISIONER                RECLAIMPOLICY   VOLUMEBINDINGMODE   ALLOWVOLUMEEXPANSION   AGE
nfs-sc (default)                   cluster.local/nfs-server-nfs-server-provisioner  Delete          Immediate            true                    75s
standard (default)                 rancher.io/local-path      Delete          WaitForFirstConsumer false                  10m
```

Stats



2. Architecture diagram of the created application



3. Component roles and connections

- **Job**
 - One-time Kubernetes pod that mounts the shared Persistent Volume Claim (PVC).
 - Writes a sample `index.html` file into the volume.
 - Ensures the web server has content to serve.
- **Deployment**
 - Manages a pod running an `nginx:alpine` container.
 - Mounts the same PVC used by the Job.
 - Serves static web content over HTTP.
- **Persistent Volume (PV)**
 - Dynamically created by the NFS provisioner in response to a PVC request.
 - Backed by the NFS server and supports `ReadWriteMany`.
 - Acts as shared storage between the Job and the NGINX pod.
- **NFS Provisioner**
 - Deployed via Helm.
 - Listens for PVC requests and dynamically provisions NFS-backed PVs.
 - Works with the custom StorageClass `nfs-sc`.
- **NFS Server**
 - Pod running NFS-Ganesha within the Kubernetes cluster.
 - Provides the physical file system backing the dynamically created PVs.
 - Shared between multiple pods via NFS protocol.

- **Kubernetes Service**

- Type `NodePort` service that exposes the NGINX pod externally.
- Allows the user to access the web app via <http://localhost:8080>.
- Routes external HTTP traffic to the correct pod inside the cluster.

Persistent Storage in Kubernetes

- By default, pod storage in Kubernetes is **ephemeral** – it disappears when the pod is deleted or restarted.

- **Persistent Volumes (PVs)** provide long-term, reliable storage that lives independently of pods.

- Applications can **claim** persistent storage by creating a **Persistent Volume Claim (PVC)**.

- In this project, persistent storage is backed by an internal **NFS server**, making it possible to share data between multiple pods.

- The **ReadWriteMany (RWX)** access mode allows the volume to be mounted by more than one pod at the same time.

- This setup ensures that:

- The **Job** can write files (e.g., `index.html`) to the volume
- The **NGINX pod** can serve those files through HTTP
 - Persistent storage is crucial for applications that:
- Need to **retain data** between pod restarts
- Share files across **multiple pods**
- Store user-generated content or configuration